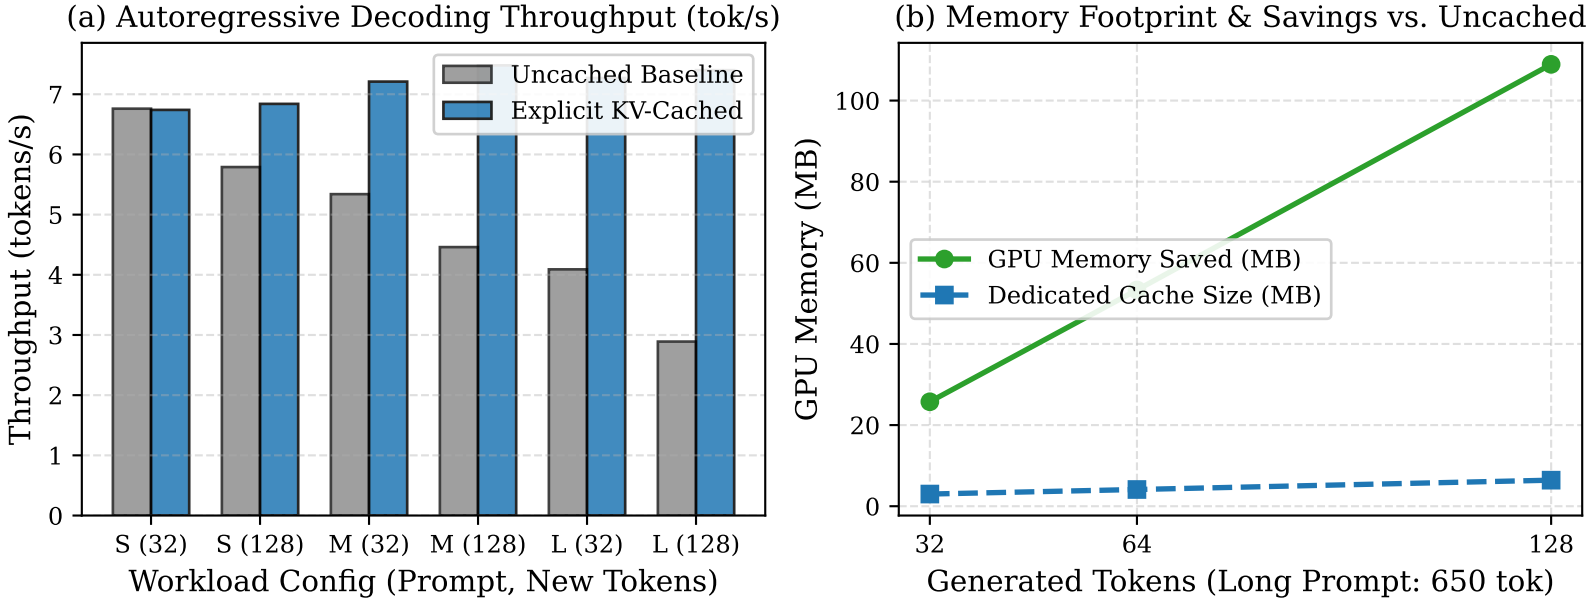
(a) Autoregressive Decoding Throughput (tok/s)
Uncached Baseline
Explicit KV-Cached
Throughput (tokens/s)
0
1
2
3
4
5
6
7
S (32)
S (128)
M (32)
M (128)
L (32)
L (128)
Workload Config (Prompt, New Tokens)
(b) Memory Footprint & Savings vs. Uncached
GPU Memory Saved (MB)
Dedicated Cache Size (MB)
GPU Memory (MB)
0
20
40
60
80
100
32
64
128
Generated Tokens (Long Prompt: 650 tok)